

## Keras -- MLPs on MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use th
is command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    plt.show()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d
, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%  
d)"%(X_train.shape[1]))  
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d  
)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normlizing  
print(X_train[0])
```

[0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

## Softmax classifier

In [11]:

```

# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```



In [12]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [13]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

In [15]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metric_s=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 45us/step - loss: 1.2891  
- accuracy: 0.6943 - val\_loss: 0.8131 - val\_accuracy: 0.8324

Epoch 2/20

60000/60000 [=====] - 3s 44us/step - loss: 0.7177  
- accuracy: 0.8414 - val\_loss: 0.6071 - val\_accuracy: 0.8615

Epoch 3/20

60000/60000 [=====] - 3s 44us/step - loss: 0.5879  
- accuracy: 0.8604 - val\_loss: 0.5251 - val\_accuracy: 0.8721

Epoch 4/20

60000/60000 [=====] - 3s 45us/step - loss: 0.5259  
- accuracy: 0.8699 - val\_loss: 0.4795 - val\_accuracy: 0.8793

Epoch 5/20

60000/60000 [=====] - 2s 32us/step - loss: 0.4883  
- accuracy: 0.8757 - val\_loss: 0.4502 - val\_accuracy: 0.8842

Epoch 6/20

60000/60000 [=====] - 2s 33us/step - loss: 0.4624  
- accuracy: 0.8801 - val\_loss: 0.4286 - val\_accuracy: 0.8883

Epoch 7/20

60000/60000 [=====] - 2s 37us/step - loss: 0.4433  
- accuracy: 0.8835 - val\_loss: 0.4128 - val\_accuracy: 0.8920

Epoch 8/20

60000/60000 [=====] - 2s 34us/step - loss: 0.4283  
- accuracy: 0.8866 - val\_loss: 0.4001 - val\_accuracy: 0.8948

Epoch 9/20

60000/60000 [=====] - 3s 46us/step - loss: 0.4163  
- accuracy: 0.8885 - val\_loss: 0.3896 - val\_accuracy: 0.8975

Epoch 10/20

60000/60000 [=====] - 2s 42us/step - loss: 0.4063  
- accuracy: 0.8908 - val\_loss: 0.3808 - val\_accuracy: 0.8992

Epoch 11/20

60000/60000 [=====] - 2s 42us/step - loss: 0.3978  
- accuracy: 0.8928 - val\_loss: 0.3735 - val\_accuracy: 0.9007

Epoch 12/20

60000/60000 [=====] - 3s 43us/step - loss: 0.3906  
- accuracy: 0.8941 - val\_loss: 0.3671 - val\_accuracy: 0.9023

Epoch 13/20

60000/60000 [=====] - 3s 49us/step - loss: 0.3841  
- accuracy: 0.8957 - val\_loss: 0.3616 - val\_accuracy: 0.9036

Epoch 14/20

60000/60000 [=====] - 3s 47us/step - loss: 0.3785  
- accuracy: 0.8968 - val\_loss: 0.3566 - val\_accuracy: 0.9048

Epoch 15/20

60000/60000 [=====] - 3s 52us/step - loss: 0.3735  
- accuracy: 0.8975 - val\_loss: 0.3523 - val\_accuracy: 0.9065

Epoch 16/20

60000/60000 [=====] - 4s 63us/step - loss: 0.3690  
- accuracy: 0.8988 - val\_loss: 0.3484 - val\_accuracy: 0.9066

Epoch 17/20

60000/60000 [=====] - 3s 45us/step - loss: 0.3649  
- accuracy: 0.8996 - val\_loss: 0.3449 - val\_accuracy: 0.9073

Epoch 18/20

60000/60000 [=====] - 3s 44us/step - loss: 0.3611  
- accuracy: 0.9005 - val\_loss: 0.3415 - val\_accuracy: 0.9086

Epoch 19/20

60000/60000 [=====] - 3s 45us/step - loss: 0.3577  
- accuracy: 0.9015 - val\_loss: 0.3386 - val\_accuracy: 0.9082

Epoch 20/20

60000/60000 [=====] - 3s 45us/step - loss: 0.3545  
- accuracy: 0.9022 - val\_loss: 0.3360 - val\_accuracy: 0.9089

In [16]:

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.33595913439393044

Test accuracy: 0.9089000225067139



# MLP + Sigmoid activation + SGDOptimizer

In [17]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
dense_2 (Dense)	(None, 512)	401920
=====		
dense_3 (Dense)	(None, 128)	65664
=====		
dense_4 (Dense)	(None, 10)	1290
=====		
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		
=====		

In [18]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 15s 244us/step - loss: 2.26

99 - accuracy: 0.2155 - val\_loss: 2.2243 - val\_accuracy: 0.3367

Epoch 2/20

60000/60000 [=====] - 15s 250us/step - loss: 2.17

93 - accuracy: 0.4564 - val\_loss: 2.1231 - val\_accuracy: 0.4867

Epoch 3/20

60000/60000 [=====] - 15s 253us/step - loss: 2.06

06 - accuracy: 0.5884 - val\_loss: 1.9769 - val\_accuracy: 0.6208

Epoch 4/20

60000/60000 [=====] - 14s 225us/step - loss: 1.88

99 - accuracy: 0.6447 - val\_loss: 1.7730 - val\_accuracy: 0.6918

Epoch 5/20

60000/60000 [=====] - 13s 216us/step - loss: 1.66

60 - accuracy: 0.6815 - val\_loss: 1.5293 - val\_accuracy: 0.7316

Epoch 6/20

60000/60000 [=====] - 13s 223us/step - loss: 1.42

72 - accuracy: 0.7132 - val\_loss: 1.2979 - val\_accuracy: 0.7497

Epoch 7/20

60000/60000 [=====] - 14s 225us/step - loss: 1.22

06 - accuracy: 0.7416 - val\_loss: 1.1150 - val\_accuracy: 0.7724

Epoch 8/20

60000/60000 [=====] - 16s 267us/step - loss: 1.06

24 - accuracy: 0.7653 - val\_loss: 0.9787 - val\_accuracy: 0.7866

Epoch 9/20

60000/60000 [=====] - 14s 225us/step - loss: 0.94

38 - accuracy: 0.7834 - val\_loss: 0.8764 - val\_accuracy: 0.7985

Epoch 10/20

60000/60000 [=====] - 15s 249us/step - loss: 0.85

30 - accuracy: 0.7976 - val\_loss: 0.7961 - val\_accuracy: 0.8093

Epoch 11/20

60000/60000 [=====] - 12s 207us/step - loss: 0.78

16 - accuracy: 0.8097 - val\_loss: 0.7324 - val\_accuracy: 0.8203

Epoch 12/20

60000/60000 [=====] - 15s 247us/step - loss: 0.72

40 - accuracy: 0.8201 - val\_loss: 0.6808 - val\_accuracy: 0.8292

Epoch 13/20

60000/60000 [=====] - 14s 234us/step - loss: 0.67

67 - accuracy: 0.8287 - val\_loss: 0.6380 - val\_accuracy: 0.8380

Epoch 14/20

60000/60000 [=====] - 14s 237us/step - loss: 0.63

74 - accuracy: 0.8363 - val\_loss: 0.6018 - val\_accuracy: 0.8461

Epoch 15/20

60000/60000 [=====] - 15s 250us/step - loss: 0.60

41 - accuracy: 0.8439 - val\_loss: 0.5717 - val\_accuracy: 0.8519

Epoch 16/20

60000/60000 [=====] - 14s 238us/step - loss: 0.57

58 - accuracy: 0.8503 - val\_loss: 0.5452 - val\_accuracy: 0.8571

Epoch 17/20

60000/60000 [=====] - 13s 215us/step - loss: 0.55

12 - accuracy: 0.8554 - val\_loss: 0.5227 - val\_accuracy: 0.8649

Epoch 18/20

60000/60000 [=====] - 14s 239us/step - loss: 0.52

98 - accuracy: 0.8609 - val\_loss: 0.5027 - val\_accuracy: 0.8681

Epoch 19/20

60000/60000 [=====] - 13s 218us/step - loss: 0.51

08 - accuracy: 0.8650 - val\_loss: 0.4850 - val\_accuracy: 0.8714

Epoch 20/20

60000/60000 [=====] - 15s 247us/step - loss: 0.49

42 - accuracy: 0.8692 - val\_loss: 0.4695 - val\_accuracy: 0.8762

In [19]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.4695337440490723

Test accuracy: 0.8762000203132629





In [20]:


```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + Sigmoid activation + ADAM

In [ ]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

In [ ]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

In [32]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

## MLP + ReLU +SGD

In [33]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4409: The name tf.random\_normal is deprecated. Please use tf.random.normal instead.

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [34]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 40us/step - loss: 0.7508

- acc: 0.7941 - val\_loss: 0.3916 - val\_acc: 0.8959

Epoch 2/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3573

- acc: 0.9000 - val\_loss: 0.3031 - val\_acc: 0.9159

Epoch 3/20

60000/60000 [=====] - 2s 37us/step - loss: 0.2923

- acc: 0.9172 - val\_loss: 0.2613 - val\_acc: 0.9273

Epoch 4/20

60000/60000 [=====] - 2s 37us/step - loss: 0.2571

- acc: 0.9273 - val\_loss: 0.2366 - val\_acc: 0.9344

Epoch 5/20

60000/60000 [=====] - 2s 35us/step - loss: 0.2325

- acc: 0.9344 - val\_loss: 0.2184 - val\_acc: 0.9387

Epoch 6/20

60000/60000 [=====] - 2s 36us/step - loss: 0.2137

- acc: 0.9398 - val\_loss: 0.2057 - val\_acc: 0.9406

Epoch 7/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1986

- acc: 0.9437 - val\_loss: 0.1926 - val\_acc: 0.9462

Epoch 8/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1858

- acc: 0.9474 - val\_loss: 0.1812 - val\_acc: 0.9472

Epoch 9/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1747

- acc: 0.9506 - val\_loss: 0.1759 - val\_acc: 0.9492

Epoch 10/20

60000/60000 [=====] - 2s 37us/step - loss: 0.1651

- acc: 0.9536 - val\_loss: 0.1646 - val\_acc: 0.9511

Epoch 11/20

60000/60000 [=====] - 2s 37us/step - loss: 0.1566

- acc: 0.9559 - val\_loss: 0.1588 - val\_acc: 0.9527

Epoch 12/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1489

- acc: 0.9582 - val\_loss: 0.1521 - val\_acc: 0.9535

Epoch 13/20

60000/60000 [=====] - 2s 34us/step - loss: 0.1421

- acc: 0.9604 - val\_loss: 0.1472 - val\_acc: 0.9553

Epoch 14/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1358

- acc: 0.9625 - val\_loss: 0.1447 - val\_acc: 0.9575

Epoch 15/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1304

- acc: 0.9640 - val\_loss: 0.1389 - val\_acc: 0.9578

Epoch 16/20

60000/60000 [=====] - 2s 37us/step - loss: 0.1252

- acc: 0.9658 - val\_loss: 0.1336 - val\_acc: 0.9592

Epoch 17/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1204

- acc: 0.9668 - val\_loss: 0.1312 - val\_acc: 0.9607

Epoch 18/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1157

- acc: 0.9679 - val\_loss: 0.1273 - val\_acc: 0.9609

Epoch 19/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1117

- acc: 0.9693 - val\_loss: 0.1242 - val\_acc: 0.9623

Epoch 20/20

60000/60000 [=====] - 2s 37us/step - loss: 0.1077

- acc: 0.9702 - val\_loss: 0.1218 - val\_acc: 0.9620

In [35]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbse=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12176313624158501

Test accuracy: 0.962

In [36]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

## MLP + ReLU + ADAM



In [37]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 46us/step - loss: 0.2240

- acc: 0.9326 - val\_loss: 0.1131 - val\_acc: 0.9679

Epoch 2/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0854

- acc: 0.9740 - val\_loss: 0.0898 - val\_acc: 0.9737

Epoch 3/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0521

- acc: 0.9842 - val\_loss: 0.0764 - val\_acc: 0.9760

Epoch 4/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0359

- acc: 0.9886 - val\_loss: 0.0726 - val\_acc: 0.9787

Epoch 5/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0278

- acc: 0.9911 - val\_loss: 0.0807 - val\_acc: 0.9757

Epoch 6/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0199

- acc: 0.9937 - val\_loss: 0.0834 - val\_acc: 0.9771

Epoch 7/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0176

- acc: 0.9947 - val\_loss: 0.0843 - val\_acc: 0.9772

Epoch 8/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0164

- acc: 0.9946 - val\_loss: 0.0779 - val\_acc: 0.9784

Epoch 9/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0117

- acc: 0.9963 - val\_loss: 0.0768 - val\_acc: 0.9799

Epoch 10/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0123

- acc: 0.9961 - val\_loss: 0.0896 - val\_acc: 0.9776

Epoch 11/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0137

- acc: 0.9954 - val\_loss: 0.0857 - val\_acc: 0.9786

Epoch 12/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0106

- acc: 0.9965 - val\_loss: 0.0854 - val\_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0081

- acc: 0.9974 - val\_loss: 0.0976 - val\_acc: 0.9775

Epoch 14/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0134

- acc: 0.9956 - val\_loss: 0.0950 - val\_acc: 0.9782

Epoch 15/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0076

- acc: 0.9977 - val\_loss: 0.0835 - val\_acc: 0.9810

```
Epoch 16/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0057
- acc: 0.9983 - val_loss: 0.1038 - val_acc: 0.9789
Epoch 17/20
60000/60000 [=====] - 2s 42us/step - loss: 0.0115
- acc: 0.9960 - val_loss: 0.0916 - val_acc: 0.9799
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0061
- acc: 0.9981 - val_loss: 0.0853 - val_acc: 0.9821
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0076
- acc: 0.9976 - val_loss: 0.1079 - val_acc: 0.9767
Epoch 20/20
60000/60000 [=====] - 2s 42us/step - loss: 0.0111
- acc: 0.9965 - val_loss: 0.0984 - val_acc: 0.9803
```

In [38]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.09837867890719448
Test accuracy: 0.9803
```

In [39]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

## MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [15]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_4 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_5 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [41]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 81us/step - loss: 0.3037

- acc: 0.9101 - val\_loss: 0.2233 - val\_acc: 0.9354

Epoch 2/20

60000/60000 [=====] - 4s 67us/step - loss: 0.1737

- acc: 0.9496 - val\_loss: 0.1665 - val\_acc: 0.9505

Epoch 3/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1356

- acc: 0.9601 - val\_loss: 0.1488 - val\_acc: 0.9546

Epoch 4/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1128

- acc: 0.9662 - val\_loss: 0.1408 - val\_acc: 0.9591

Epoch 5/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0959

- acc: 0.9713 - val\_loss: 0.1229 - val\_acc: 0.9624

Epoch 6/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0821

- acc: 0.9747 - val\_loss: 0.1190 - val\_acc: 0.9634

Epoch 7/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0681

- acc: 0.9781 - val\_loss: 0.1126 - val\_acc: 0.9652

Epoch 8/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0599

- acc: 0.9819 - val\_loss: 0.1055 - val\_acc: 0.9681

Epoch 9/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0520

- acc: 0.9833 - val\_loss: 0.1084 - val\_acc: 0.9668

Epoch 10/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0456

- acc: 0.9852 - val\_loss: 0.1095 - val\_acc: 0.9677

Epoch 11/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0389

- acc: 0.9874 - val\_loss: 0.0959 - val\_acc: 0.9724

Epoch 12/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0347

- acc: 0.9888 - val\_loss: 0.0942 - val\_acc: 0.9727

Epoch 13/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0308

- acc: 0.9899 - val\_loss: 0.0996 - val\_acc: 0.9712

Epoch 14/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0297

- acc: 0.9900 - val\_loss: 0.1011 - val\_acc: 0.9702

Epoch 15/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0253

- acc: 0.9919 - val\_loss: 0.0966 - val\_acc: 0.9746

Epoch 16/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0202

- acc: 0.9937 - val\_loss: 0.0874 - val\_acc: 0.9751

Epoch 17/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0216

- acc: 0.9930 - val\_loss: 0.0954 - val\_acc: 0.9723

Epoch 18/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0199

- acc: 0.9935 - val\_loss: 0.0985 - val\_acc: 0.9723

Epoch 19/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0182

- acc: 0.9943 - val\_loss: 0.0912 - val\_acc: 0.9743

Epoch 20/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0182

- acc: 0.9937 - val\_loss: 0.0998 - val\_acc: 0.9710

In [0]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09344355644605822

Test accuracy: 0.9757



In [0]:

```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

## 5. MLP + Dropout + AdamOptimizer

In [42]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [43]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 86us/step - loss: 0.6731  
- acc: 0.7920 - val\_loss: 0.2850 - val\_acc: 0.9165

Epoch 2/20

60000/60000 [=====] - 4s 71us/step - loss: 0.4320  
- acc: 0.8679 - val\_loss: 0.2531 - val\_acc: 0.9277

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.3834  
- acc: 0.8831 - val\_loss: 0.2284 - val\_acc: 0.9334

Epoch 4/20

60000/60000 [=====] - 4s 74us/step - loss: 0.3546  
- acc: 0.8919 - val\_loss: 0.2173 - val\_acc: 0.9347

Epoch 5/20

60000/60000 [=====] - 4s 71us/step - loss: 0.3393  
- acc: 0.8976 - val\_loss: 0.2106 - val\_acc: 0.9384

Epoch 6/20

60000/60000 [=====] - 4s 72us/step - loss: 0.3192  
- acc: 0.9044 - val\_loss: 0.1996 - val\_acc: 0.9391

Epoch 7/20

60000/60000 [=====] - 4s 69us/step - loss: 0.3063  
- acc: 0.9067 - val\_loss: 0.1914 - val\_acc: 0.9434

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2941  
- acc: 0.9108 - val\_loss: 0.1823 - val\_acc: 0.9464

Epoch 9/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2825  
- acc: 0.9151 - val\_loss: 0.1755 - val\_acc: 0.9470

Epoch 10/20

60000/60000 [=====] - 4s 73us/step - loss: 0.2682  
- acc: 0.9200 - val\_loss: 0.1653 - val\_acc: 0.9512

Epoch 11/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2560  
- acc: 0.9234 - val\_loss: 0.1593 - val\_acc: 0.9518

Epoch 12/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2457  
- acc: 0.9259 - val\_loss: 0.1532 - val\_acc: 0.9545

Epoch 13/20

60000/60000 [=====] - 4s 72us/step - loss: 0.2367  
- acc: 0.9290 - val\_loss: 0.1488 - val\_acc: 0.9557

Epoch 14/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2264  
- acc: 0.9314 - val\_loss: 0.1350 - val\_acc: 0.9593

Epoch 15/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2182  
- acc: 0.9328 - val\_loss: 0.1305 - val\_acc: 0.9624

Epoch 16/20

60000/60000 [=====] - 4s 74us/step - loss: 0.2073  
- acc: 0.9382 - val\_loss: 0.1286 - val\_acc: 0.9619

Epoch 17/20

60000/60000 [=====] - 4s 70us/step - loss: 0.1986  
- acc: 0.9389 - val\_loss: 0.1192 - val\_acc: 0.9661

Epoch 18/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1907  
- acc: 0.9424 - val\_loss: 0.1152 - val\_acc: 0.9663

Epoch 19/20

60000/60000 [=====] - 4s 73us/step - loss: 0.1820  
- acc: 0.9459 - val\_loss: 0.1188 - val\_acc: 0.9663

Epoch 20/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1789  
- acc: 0.9463 - val\_loss: 0.1063 - val\_acc: 0.9693

In [44]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10633013175930828

Test accuracy: 0.9693

In [45]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

## Hyper-parameter tuning of Keras models using Sklearn

In [0]:

```

from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer
=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='ada
m')

    return model

```

In [47]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid', 'relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1
978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.
warnings.warn(CV_WARNING, FutureWarning)
```

In [48]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.975333 using {'activ': 'sigmoid'}
0.975333 (0.001211) with: {'activ': 'sigmoid'}
0.975083 (0.001109) with: {'activ': 'relu'}
```

In [0]:

## ASSIGNMENT

**\*\* 2 LAYER ARCHITECTURE WITH BATCH NORMALIZATION AND DROPOUT\*\***

In [16]:

```
from keras.layers import Dropout

model_2= Sequential()

model_2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_2.add(BatchNormalization())
model_2.add(Dropout(0.5))

model_2.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )
model_2.add(BatchNormalization())
model_2.add(Dropout(0.5))

model_2.add(Dense(output_dim, activation='softmax'))

model_2.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 364)	285740
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
dropout_1 (Dropout)	(None, 364)	0
dense_7 (Dense)	(None, 256)	93440
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dropout_2 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 10)	2570
Total params: 384,230		
Trainable params: 382,990		
Non-trainable params: 1,240		



In [ ]:

```
model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 124us/step - loss: 0.491

8 - accuracy: 0.8500 - val\_loss: 0.1702 - val\_accuracy: 0.9486

Epoch 2/20

60000/60000 [=====] - 6s 101us/step - loss: 0.253

5 - accuracy: 0.9239 - val\_loss: 0.1308 - val\_accuracy: 0.9586

Epoch 3/20

60000/60000 [=====] - 7s 115us/step - loss: 0.206

6 - accuracy: 0.9378 - val\_loss: 0.1118 - val\_accuracy: 0.9644

Epoch 4/20

60000/60000 [=====] - 7s 121us/step - loss: 0.170

4 - accuracy: 0.9480 - val\_loss: 0.0958 - val\_accuracy: 0.9698

Epoch 5/20

30976/60000 [=====&gt;.....] - ETA: 3s - loss: 0.1586 - ac

curacy: 0.9510

In [51]:

```
score = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.054869781568923644

Test accuracy: 0.9822

In [1]:

```
import matplotlib.pyplot as plt
x=[1,2,3,4,5]
y=[1,2,9,16,25]
plt.plot(x,y)
plt.xlabel('x')
plt.ylabel('y')
plt.grid()
plt.show()
```



In [53]:

```
w_after = model_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

**\*\* 2 LAYER ARCHITECTURE WITHOUT BATCH NORMALIZATION AND DROPOUT\*\***

In [14]:

```
from keras.layers import Dropout

model_3= Sequential()

model_3.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_3.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_2 (Dense)	(None, 364)	285740
dense_3 (Dense)	(None, 256)	93440
dense_4 (Dense)	(None, 10)	2570
=====	=====	=====
Total params: 381,750		
Trainable params: 381,750		
Non-trainable params: 0		

In [ ]:

```
model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'
])

history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
=1, validation_data=(X_test, Y_test))
```

In [56]:

```
score = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11713699789448274

Test accuracy: 0.9784

In [57]:

```
w_after = model_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

### 3 LAYER ARCHITECTURE

#### WITH BATCH NORMALIZATION AND DROPOUT

In [59]:

```
from keras.layers import Dropout

model_4= Sequential()

model_4.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_4.add(BatchNormalization())
model_4.add(Dropout(0.2))

model_4.add(Dense(356, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )
model_4.add(BatchNormalization())
model_4.add(Dropout(0.2))

model_4.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )
model_4.add(BatchNormalization())
model_4.add(Dropout(0.2))

model_4.add(Dense(output_dim, activation='softmax'))

model_4.summary()
```

Model: "sequential\_18"

Layer (type)	Output Shape	Param #
dense_50 (Dense)	(None, 512)	401920
batch_normalization_9 (Batch Normalization)	(None, 512)	2048
dropout_7 (Dropout)	(None, 512)	0
dense_51 (Dense)	(None, 356)	182628
batch_normalization_10 (Batch Normalization)	(None, 356)	1424
dropout_8 (Dropout)	(None, 356)	0
dense_52 (Dense)	(None, 256)	91392
batch_normalization_11 (Batch Normalization)	(None, 256)	1024
dropout_9 (Dropout)	(None, 256)	0
dense_53 (Dense)	(None, 10)	2570
Total params: 683,006		
Trainable params: 680,758		
Non-trainable params: 2,248		

In [60]:

```
model_4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 125us/step - loss: 0.318  
2 - acc: 0.9031 - val\_loss: 0.1210 - val\_acc: 0.9598

Epoch 2/20

60000/60000 [=====] - 5s 91us/step - loss: 0.1461  
- acc: 0.9549 - val\_loss: 0.0939 - val\_acc: 0.9704

Epoch 3/20

60000/60000 [=====] - 5s 87us/step - loss: 0.1101  
- acc: 0.9653 - val\_loss: 0.0824 - val\_acc: 0.9733

Epoch 4/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0921  
- acc: 0.9712 - val\_loss: 0.0862 - val\_acc: 0.9721

Epoch 5/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0780  
- acc: 0.9751 - val\_loss: 0.0747 - val\_acc: 0.9766

Epoch 6/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0712  
- acc: 0.9777 - val\_loss: 0.0662 - val\_acc: 0.9784

Epoch 7/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0651  
- acc: 0.9790 - val\_loss: 0.0697 - val\_acc: 0.9792

Epoch 8/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0552  
- acc: 0.9818 - val\_loss: 0.0711 - val\_acc: 0.9789

Epoch 9/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0524  
- acc: 0.9824 - val\_loss: 0.0695 - val\_acc: 0.9801

Epoch 10/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0502  
- acc: 0.9834 - val\_loss: 0.0670 - val\_acc: 0.9820

Epoch 11/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0459  
- acc: 0.9848 - val\_loss: 0.0631 - val\_acc: 0.9812

Epoch 12/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0443  
- acc: 0.9855 - val\_loss: 0.0694 - val\_acc: 0.9789

Epoch 13/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0379  
- acc: 0.9874 - val\_loss: 0.0643 - val\_acc: 0.9822

Epoch 14/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0354  
- acc: 0.9879 - val\_loss: 0.0655 - val\_acc: 0.9828

Epoch 15/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0327  
- acc: 0.9894 - val\_loss: 0.0696 - val\_acc: 0.9807

Epoch 16/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0337  
- acc: 0.9889 - val\_loss: 0.0637 - val\_acc: 0.9838

Epoch 17/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0298  
- acc: 0.9902 - val\_loss: 0.0675 - val\_acc: 0.9812

Epoch 18/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0323  
- acc: 0.9892 - val\_loss: 0.0684 - val\_acc: 0.9817

Epoch 19/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0278  
- acc: 0.9904 - val\_loss: 0.0727 - val\_acc: 0.9806

Epoch 20/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0264  
- acc: 0.9911 - val\_loss: 0.0653 - val\_acc: 0.9826



In [61]:

```
score = model_4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06527992642150567

Test accuracy: 0.9826

In [62]:

```
print(score)
```

[0.06527992642150567, 0.9826]

In [64]:

```
w_after = model_4.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='g')
plt.xlabel('OUTPUT LAYER WEIGHTS ')
plt.show()
```

## WITHOUT BATCH NORMALIZATION AND DROPOUT

In [69]:

```

from keras.layers import Dropout

model_5= Sequential()

model_5.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_5.add(Dense(356, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )

model_5.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )

model_5.add(Dense(output_dim, activation='softmax'))

model_5.summary()

```

Model: "sequential\_19"

Layer (type)	Output Shape	Param #
=====		
dense_54 (Dense)	(None, 512)	401920
dense_55 (Dense)	(None, 356)	182628
dense_56 (Dense)	(None, 256)	91392
dense_57 (Dense)	(None, 10)	2570
=====		
Total params: 678,510		
Trainable params: 678,510		
Non-trainable params: 0		
=====		

In [70]:

```
model_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 81us/step - loss: 3.6271

- acc: 0.7508 - val\_loss: 2.2213 - val\_acc: 0.8383

Epoch 2/20

60000/60000 [=====] - 3s 48us/step - loss: 2.1185

- acc: 0.8489 - val\_loss: 2.0017 - val\_acc: 0.8566

Epoch 3/20

60000/60000 [=====] - 3s 48us/step - loss: 1.9423

- acc: 0.8620 - val\_loss: 1.8992 - val\_acc: 0.8636

Epoch 4/20

60000/60000 [=====] - 3s 50us/step - loss: 1.8517

- acc: 0.8703 - val\_loss: 1.8494 - val\_acc: 0.8682

Epoch 5/20

60000/60000 [=====] - 3s 50us/step - loss: 1.8119

- acc: 0.8738 - val\_loss: 1.8758 - val\_acc: 0.8637

Epoch 6/20

60000/60000 [=====] - 3s 51us/step - loss: 1.7877

- acc: 0.8766 - val\_loss: 1.8339 - val\_acc: 0.8695

Epoch 7/20

60000/60000 [=====] - 3s 51us/step - loss: 0.7151

- acc: 0.9378 - val\_loss: 0.2033 - val\_acc: 0.9661

Epoch 8/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1027

- acc: 0.9800 - val\_loss: 0.2098 - val\_acc: 0.9661

Epoch 9/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0742

- acc: 0.9838 - val\_loss: 0.1661 - val\_acc: 0.9709

Epoch 10/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0688

- acc: 0.9850 - val\_loss: 0.1831 - val\_acc: 0.9693

Epoch 11/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0610

- acc: 0.9862 - val\_loss: 0.1754 - val\_acc: 0.9681

Epoch 12/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0578

- acc: 0.9866 - val\_loss: 0.1677 - val\_acc: 0.9692

Epoch 13/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0531

- acc: 0.9873 - val\_loss: 0.1552 - val\_acc: 0.9723

Epoch 14/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0492

- acc: 0.9890 - val\_loss: 0.1644 - val\_acc: 0.9737

Epoch 15/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0581

- acc: 0.9869 - val\_loss: 0.1597 - val\_acc: 0.9724

Epoch 16/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0413

- acc: 0.9904 - val\_loss: 0.1712 - val\_acc: 0.9712

Epoch 17/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0421

- acc: 0.9904 - val\_loss: 0.1707 - val\_acc: 0.9725

Epoch 18/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0393

- acc: 0.9911 - val\_loss: 0.1690 - val\_acc: 0.9736

Epoch 19/20

60000/60000 [=====] - 3s 48us/step - loss: 0.0393

- acc: 0.9910 - val\_loss: 0.1581 - val\_acc: 0.9739

Epoch 20/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0430

- acc: 0.9905 - val\_loss: 0.1619 - val\_acc: 0.9734

In [71]:

```
score = model_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.16188472667328224

Test accuracy: 0.9734

In [72]:

```
w_after = model_5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='g')
plt.xlabel('OUTPUT LAYER WEIGHTS ')
plt.show()
```

## 5 LAYER ARCHITECTURE

### WITH BATCH NORMALIZATION AND DROPOUT

In [17]:

```
from keras.layers import Dropout

model_6= Sequential()

model_6.add(Dense(656, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_6.add(BatchNormalization())
model_6.add(Dropout(0.2))

model_6.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_6.add(BatchNormalization())
model_6.add(Dropout(0.2))

model_6.add(Dense(356, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )
model_6.add(BatchNormalization())
model_6.add(Dropout(0.2))

model_6.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )
model_6.add(BatchNormalization())
model_6.add(Dropout(0.2))

model_6.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_6.add(BatchNormalization())
model_6.add(Dropout(0.2))

model_6.add(Dense(output_dim, activation='softmax'))

model_6.summary()
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 656)	514960
batch_normalization_3 (Batch Normalization)	(None, 656)	2624
dropout_1 (Dropout)	(None, 656)	0
dense_7 (Dense)	(None, 512)	336384
batch_normalization_4 (Batch Normalization)	(None, 512)	2048
dropout_2 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 356)	182628
batch_normalization_5 (Batch Normalization)	(None, 356)	1424
dropout_3 (Dropout)	(None, 356)	0
dense_9 (Dense)	(None, 256)	91392
batch_normalization_6 (Batch Normalization)	(None, 256)	1024
dropout_4 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 128)	32896
batch_normalization_7 (Batch Normalization)	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 10)	1290
Total params: 1,167,182		
Trainable params: 1,163,366		
Non-trainable params: 3,816		

In [18]:

```
model_6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/math\_grad.py:1424: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1033: The name tf.assign\_add is deprecated. Please use tf.compat.v1.assign\_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:190: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:207: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:216: The name tf.is\_variable\_initialized is deprecated. Please use tf.compat.v1.is\_variable\_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:223: The name tf.variables\_initializer is deprecated. Please use tf.compat.v1.variables\_initializer instead.

60000/60000 [=====] - 16s 271us/step - loss: 0.3319 - acc: 0.8975 - val\_loss: 0.1201 - val\_acc: 0.9633

Epoch 2/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1423 - acc: 0.9556 - val\_loss: 0.0979 - val\_acc: 0.9701

Epoch 3/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1083 - acc: 0.9666 - val\_loss: 0.0835 - val\_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0906 - acc: 0.9718 - val\_loss: 0.0749 - val\_acc: 0.9785

Epoch 5/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0808 - acc: 0.9755 - val\_loss: 0.0732 - val\_acc: 0.9795

```
Epoch 6/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0705
- acc: 0.9783 - val_loss: 0.0689 - val_acc: 0.9792
Epoch 7/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0650
- acc: 0.9795 - val_loss: 0.0713 - val_acc: 0.9809
Epoch 8/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0602
- acc: 0.9809 - val_loss: 0.0642 - val_acc: 0.9820
Epoch 9/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0533
- acc: 0.9829 - val_loss: 0.0747 - val_acc: 0.9797
Epoch 10/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0462
- acc: 0.9848 - val_loss: 0.0677 - val_acc: 0.9811
Epoch 11/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0461
- acc: 0.9857 - val_loss: 0.0576 - val_acc: 0.9830
Epoch 12/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0459
- acc: 0.9851 - val_loss: 0.0621 - val_acc: 0.9824
Epoch 13/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0405
- acc: 0.9872 - val_loss: 0.0643 - val_acc: 0.9821
Epoch 14/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0366
- acc: 0.9882 - val_loss: 0.0656 - val_acc: 0.9814
Epoch 15/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0352
- acc: 0.9887 - val_loss: 0.0642 - val_acc: 0.9821
Epoch 16/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0355
- acc: 0.9884 - val_loss: 0.0647 - val_acc: 0.9838
Epoch 17/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0327
- acc: 0.9897 - val_loss: 0.0622 - val_acc: 0.9835
Epoch 18/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0313
- acc: 0.9894 - val_loss: 0.0642 - val_acc: 0.9832
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0298
- acc: 0.9901 - val_loss: 0.0585 - val_acc: 0.9844
Epoch 20/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0260
- acc: 0.9917 - val_loss: 0.0605 - val_acc: 0.9848
```

In [19]:

```
score = model_6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06047782948987442

Test accuracy: 0.9848

In [22]:

```
w_after = model_6.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w =w_after[4].flatten().reshape(-1,1)
h4_w =w_after[6].flatten().reshape(-1,1)
h5_w =w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2,3,3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='g')
plt.xlabel('OUTPUT LAYER WEIGHTS ')
plt.show()
```

## WITHOUT BATCH NORMALIZATION AND DROPOUT

In [23]:

```
from keras.layers import Dropout

model_7= Sequential()

model_7.add(Dense(656, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_7.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_7.add(Dense(356, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )

model_7.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.55, seed=None)) )

model_7.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=
RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_7.add(Dense(output_dim, activation='softmax'))

model_7.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
dense_12 (Dense)	(None, 656)	514960
dense_13 (Dense)	(None, 512)	336384
dense_14 (Dense)	(None, 356)	182628
dense_15 (Dense)	(None, 256)	91392
dense_16 (Dense)	(None, 128)	32896
dense_17 (Dense)	(None, 10)	1290
=====		
Total params: 1,159,550		
Trainable params: 1,159,550		
Non-trainable params: 0		
=====		

In [24]:

```
model_7.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_7.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 54us/step - loss: 1.7796  
- acc: 0.8432 - val\_loss: 0.2184 - val\_acc: 0.9336

Epoch 2/20

60000/60000 [=====] - 3s 42us/step - loss: 0.1299  
- acc: 0.9609 - val\_loss: 0.1096 - val\_acc: 0.9680

Epoch 3/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0793  
- acc: 0.9756 - val\_loss: 0.0978 - val\_acc: 0.9703

Epoch 4/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0637  
- acc: 0.9805 - val\_loss: 0.0898 - val\_acc: 0.9760

Epoch 5/20

60000/60000 [=====] - 2s 42us/step - loss: 0.0517  
- acc: 0.9836 - val\_loss: 0.0916 - val\_acc: 0.9742

Epoch 6/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0478  
- acc: 0.9845 - val\_loss: 0.0996 - val\_acc: 0.9735

Epoch 7/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0430  
- acc: 0.9868 - val\_loss: 0.0860 - val\_acc: 0.9774

Epoch 8/20

60000/60000 [=====] - 3s 44us/step - loss: 0.0335  
- acc: 0.9899 - val\_loss: 0.1082 - val\_acc: 0.9732

Epoch 9/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0359  
- acc: 0.9887 - val\_loss: 0.1383 - val\_acc: 0.9680

Epoch 10/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0325  
- acc: 0.9904 - val\_loss: 0.1018 - val\_acc: 0.9777

Epoch 11/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0308  
- acc: 0.9905 - val\_loss: 0.1037 - val\_acc: 0.9743

Epoch 12/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0282  
- acc: 0.9912 - val\_loss: 0.0858 - val\_acc: 0.9788

Epoch 13/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0249  
- acc: 0.9924 - val\_loss: 0.0953 - val\_acc: 0.9792

Epoch 14/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0249  
- acc: 0.9922 - val\_loss: 0.1075 - val\_acc: 0.9772

Epoch 15/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0199  
- acc: 0.9943 - val\_loss: 0.1078 - val\_acc: 0.9783

Epoch 16/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0227  
- acc: 0.9929 - val\_loss: 0.1000 - val\_acc: 0.9752

Epoch 17/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0214  
- acc: 0.9939 - val\_loss: 0.0980 - val\_acc: 0.9798

Epoch 18/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0150  
- acc: 0.9953 - val\_loss: 0.1106 - val\_acc: 0.9778

Epoch 19/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0193  
- acc: 0.9950 - val\_loss: 0.1143 - val\_acc: 0.9781

Epoch 20/20

60000/60000 [=====] - 2s 42us/step - loss: 0.0199  
- acc: 0.9944 - val\_loss: 0.1074 - val\_acc: 0.9790

In [25]:

```
score = model_7.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10735029972919778

Test accuracy: 0.979

In [26]:

```
w_after = model_7.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w =w_after[4].flatten().reshape(-1,1)
h4_w =w_after[6].flatten().reshape(-1,1)
h5_w =w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='b')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2,3,3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='g')
plt.xlabel('OUTPUT LAYER WEIGHTS ')
plt.show()
```

## CONCLUSION

In [27]:

```

from prettytable import PrettyTable
x=PrettyTable(['ACTIVATION','OPTIMIZER','NO OF HIDDEN LAYERS','BATCHNORM AND DROPOUT',
'TEST ACCURACY','TEST SCORE'])
x.add_row(['RELU','ADAM','2','YES','0.9822','0.0548'])
x.add_row(['RELU','ADAM','2','NO','0.9784','0.117136'])
x.add_row(['RELU','ADAM','3','YES','0.9826','0.06527'])
x.add_row(['RELU','ADAM','3','NO','0.9734','0.16188'])
x.add_row(['RELU','ADAM','5','YES','0.9848','0.060477'])
x.add_row(['RELU','ADAM','5','NO','0.979','0.1073'])

print(x.get_string(start=0,end=9))

```

```

+-----+-----+-----+-----+-----+
| ACTIVATION | OPTIMIZER | NO OF HIDDEN LAYERS | BATCHNORM AND DROPOUT | T
EST ACCURACY | TEST SCORE |
+-----+-----+-----+-----+-----+
| RELU      | ADAM      | 2      | YES      | 0.9822 | 0.0548 |
| RELU      | ADAM      | 2      | NO       | 0.9784 | 0.117136 |
| RELU      | ADAM      | 3      | YES      | 0.9826 | 0.06527 |
| RELU      | ADAM      | 3      | NO       | 0.9734 | 0.16188 |
| RELU      | ADAM      | 5      | YES      | 0.9848 | 0.060477 |
| RELU      | ADAM      | 5      | NO       | 0.979  | 0.1073  |
+-----+-----+-----+-----+-----+

```

In [0]:

In [0]:

In [0]: